

MADS-ML – Machine Learning

Neighborhood-based Classification

Prof. Dr. Stephan Doerfel



FACHHOCHSCHULE KIEL
University of Applied Sciences



Moodle (WiSe 2024/25)

For a decision, remember similar instances and their decisions.

Outline

Basic Idea

The Neighborhood Size

The Decision Rule

Distance Functions

Python

Basic Idea 1/2

Assumption:

Basic Idea 1/2

Assumption:

Instances of the same class are **similar**. Similar instances are of the same class.

Basic Idea 1/2

Assumption:

Instances of the same class are **similar**. Similar instances are of the same class.

Idea:

Basic Idea 1/2

Assumption:

Instances of the same class are **similar**. Similar instances are of the same class.

Idea:

For a data instance x find those k **instances** from the training data that are the **most similar** to x .

Basic Idea 1/2

Assumption:

Instances of the same class are **similar**. Similar instances are of the same class.

Idea:

For a data instance x find those k **instances** from the training data that are the **most similar** to x .

Classification:

Basic Idea 1/2

Assumption:

Instances of the same class are **similar**. Similar instances are of the same class.

Idea:

For a data instance x find those k **instances** from the training data that are the **most similar** to x .

Classification:

Derive the class of x from the classes of these k instances.



Which ingredients do we need?

Basic Idea 2/2

“Lazy Learner”:

- ▶ (as opposed to Eager Learners)
- ▶ no explicit model is trained.
- ▶ the training data serves as the model
- ▶ no generalization from the training data

Basic Idea 2/2

“Lazy Learner”:

- ▶ (as opposed to Eager Learners)
- ▶ no explicit model is trained.
- ▶ the training data serves as the model
- ▶ no generalization from the training data

Distance-based:

distance function determines the (non-)similarity between to data instances

Basic Idea 2/2

“Lazy Learner”:

- ▶ (as opposed to Eager Learners)
- ▶ no explicit model is trained.
- ▶ the training data serves as the model
- ▶ no generalization from the training data

Distance-based:


distance function determines the (non-)similarity between to data instances

How many neighbors:

select k nearest (most similar) neighbors

→ Parameter $k \Rightarrow$ “ k -nearest neighbor”-algorithm

kNN: Example




x new data instance

○ Neighborhood for $k = 1$

○ Neighborhood for $k = 7$

- ▶ training data: 11 instances of class “+” and 8 instances of class “-”.
- ▶ $k = 1$: object x is classified
- ▶ $k = 7$ objekt x is classified

kNN: Example



x new data instance

○ Neighborhood for $k = 1$

○ Neighborhood for $k = 7$

- ▶ training data: 11 instances of class "+" and 8 instances of class "-".
- ▶ $k = 1$: object x is classified "+"
- ▶ $k = 7$ objekt x is classified "-"

kNN algorithm

Given a data instance x

- ▶ for each instance t of the training data
 - ▶ compute distance between x and t
- ▶ select k instances with the lowest distance
- ▶ apply the decision rule to those k instances

kNN Ingredients

Decision Data:

- ▶ the data set from which we choose the k nearest neighbors
- ▶ these are exactly the training data

kNN Ingredients

Decision Data:

- ▶ the data set from which we choose the k nearest neighbors
- ▶ these are exactly the training data

Neighborhood Size:

- ▶ hyperparameter k , the number of neighbors considered per instance

kNN Ingredients

Decision Data:

- ▶ the data set from which we choose the k nearest neighbors
- ▶ these are exactly the training data

Neighborhood Size:

- ▶ hyperparameter k , the number of neighbors considered per instance

Distance Function:

- ▶ defines the (non)-similarity for pairs of data instances
- ▶ many different choices depending on the data

kNN Ingredients

Decision Data:

- ▶ the data set from which we choose the k nearest neighbors
- ▶ these are exactly the training data

Neighborhood Size:

- ▶ hyperparameter k , the number of neighbors considered per instance

Distance Function:

- ▶ defines the (non)-similarity for pairs of data instances
- ▶ many different choices depending on the data

Decision Rule:

- ▶ How does one derive the class from the given k neighbors?

Outline

Basic Idea

The Neighborhood Size


The Decision Rule

Distance Functions

Python

Choice of k

- ▶ k “too low”:




$\cdots k = 1$

$\cdots k = 7$

$\bigcirc k > 7$

Choice of k

- ▶ k “too low”: high sensitivity for outliers




$\text{dotted circle} \quad k = 1$

$\text{dashed circle} \quad k = 7$

$\text{solid circle} \quad k > 7$

Choice of k

- ▶ k “too low”: high sensitivity for outliers
- ▶ k “too high”:




$\circ \quad k = 1$

$\circ \quad k = 7$

$\circ \quad k > 7$

Choice of k

- ▶ k “too low”: high sensitivity for outliers
- ▶ k “too high”: many objects from other classes in the neighborhood




$\text{dotted circle} \quad k = 1$

$\text{dashed circle} \quad k = 7$

$\text{solid circle} \quad k > 7$

Choice of k

- ▶ k “too low”: high sensitivity for outliers
- ▶ k “too high”: many objects from other classes in the neighborhood
- ▶ rule of thumb: often $1 \ll k < 10$ yields highest classification quality



$\cdots k = 1$

$- - - k = 7$

$\bigcirc k > 7$

Simple determination of k

- ▶ Try different values $k = 1, 2, 3, \dots$
- ▶ For each k : compute classification quality
 - ▶ ⌚ usually using cross validation
- ▶ use best k in production

Outline

Basic Idea

The Neighborhood Size

The Decision Rule

Distance Functions

Python

Decision rule

Decision rule

Standard:

- ▶ Use the majority class among the neighbors. Each neighbor has one vote with the weight 1.

Decision rule

Standard:

- ▶ Use the majority class among the neighbors. Each neighbor has one vote with the weight 1.

Weighted:

- ▶ Each neighbor votes with a weight:
 - ▶ by inverted distance
 - ▶ by inverted class frequency in the training dataset
 - ▶ by some cost function (e.g. when predicting $+$ instead of actual $-$ is more expensive than $-$ instead of $+$).

Decision rule

Standard:

- ▶ Use the majority class among the neighbors. Each neighbor has one vote with the weight 1.

Weighted:

- ▶ Each neighbor votes with a weight:
 - ▶ by inverted distance
 - ▶ by inverted class frequency in the training dataset
 - ▶ by some cost function (e.g. when predicting $+$ instead of actual $-$ is more expensive than $-$ instead of $+$).

Example weighted voting

Let the training set contain 95% of class A and only 5% of class B .
A neighborhood of

$$\{A, A, A, A, B, B, B\}$$

yields: Standard $\Rightarrow A$, weighted by class $\Rightarrow B$.

kNN with probabilities

- ▶ probabilities help us judge “how sure” the classifier is in its classifications

kNN with probabilities

- ▶ probabilities help us judge “how sure” the classifier is in its classifications
- ▶ kNN offers probabilities based on the voting.

kNN with probabilities

- ▶ probabilities help us judge “how sure” the classifier is in its classifications
- ▶ kNN offers probabilities based on the voting.
- ▶ add the votes per class and divide by the number of neighbors

Example probabilities

Let the training set contain instances of classes A , B , and C . A neighborhood of

$$\{A, A, A, A, B, B, B\}$$

yields probabilities $p_A(x) = \frac{4}{7}$, $p_B(x) = \frac{3}{7}$, and $p_C(x) = 0$.

Outline

Basic Idea


The Neighborhood Size

The Decision Rule

Distance Functions

Python

Distance functions


Just a couple of examples. More on distances in  MADS-MMS.

Definition 1 (Minkowski-Metrik)

Let $x = (x_1, \dots, x_d)$ and $y = (y_1, \dots, y_d)$ be numerical vectors. Then the **Minkowski-Metric** for p (L_p -metric) is defined as:

$$\text{dist}_p(x, y) := \sqrt[p]{\sum_{i=1}^d |x_i - y_i|^p}$$

Distance functions

Just a couple of examples. More on distances in  MADS-MMS.

Definition 1 (Minkowski-Metrik)

Let $x = (x_1, \dots, x_d)$ and $y = (y_1, \dots, y_d)$ be numerical vectors. Then the **Minkowski-Metric** for p (L_p -metric) is defined as:

$$\text{dist}_p(x, y) := \sqrt[p]{\sum_{i=1}^d |x_i - y_i|^p}$$

Commonly used Minkowski-Metrics are:

- ▶ Euclidean Distance ($p = 2$): $\text{dist}_2(x, y) = \sqrt{\sum_{i=1}^d |x_i - y_i|^2}$
- ▶ Manhattan-Metric ($p = 1$): $\text{dist}_1(x, y) = \sum_{i=1}^d |x_i - y_i|$
- ▶ Maximum-Metric ($p = \infty$):
 $\text{dist}_\infty(x, y) = \max\{|x_i - y_i| \mid 1 \leq i \leq d\}$

A Problem with Distances

Imagine a dataset where the data instances are employees and the features are age, yearly salary and number of children. Using the Euclidean distance we yield:

$$\text{dist}_2(x, y) = \sqrt{(x_{\text{age}} - y_{\text{age}})^2 + (x_{\text{kids}} - y_{\text{kids}})^2 + (x_{\text{salary}} - y_{\text{salary}})^2}$$

A Problem with Distances

Imagine a dataset where the data instances are employees and the features are age, yearly salary and number of children. Using the Euclidean distance we yield:

$$\text{dist}_2(x, y) = \sqrt{(x_{\text{age}} - y_{\text{age}})^2 + (x_{\text{kids}} - y_{\text{kids}})^2 + (x_{\text{salary}} - y_{\text{salary}})^2}$$

- ▶ A difference of 1 is a lot for the number of kids, it is very little for the yearly salary.

A Problem with Distances

Imagine a dataset where the data instances are employees and the features are age, yearly salary and number of children. Using the Euclidean distance we yield:

$$\text{dist}_2(x, y) = \sqrt{(x_{\text{age}} - y_{\text{age}})^2 + (x_{\text{kids}} - y_{\text{kids}})^2 + (x_{\text{salary}} - y_{\text{salary}})^2}$$

- ▶ A difference of 1 is a lot for the number of kids, it is very little for the yearly salary.
- ▶ salary dominates the other features

A Problem with Distances

Imagine a dataset where the data instances are employees and the features are age, yearly salary and number of children. Using the Euclidean distance we yield:

$$\text{dist}_2(x, y) = \sqrt{(x_{\text{age}} - y_{\text{age}})^2 + (x_{\text{kids}} - y_{\text{kids}})^2 + (x_{\text{salary}} - y_{\text{salary}})^2}$$

- ▶ A difference of 1 is a lot for the number of kids, it is very little for the yearly salary.
- ▶ salary dominates the other features
- ▶ solution: **feature scaling** (normalization)

Feature Scaling (Min-Max)

Definition 2

Min-Max-Scaling a feature in a dataset D means computing

$$x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

where

- ▶ x is the original values of that feature for some data instance
 - ▶ x' is the scaled value of x , and
 - ▶ x_{\min} and x_{\max} are the lowest and highest values resp. in the dataset.
-
- ▶ Scaling maps the values of a feature into the interval

Feature Scaling (Min-Max)

Definition 2

Min-Max-Scaling a feature in a dataset D means computing

$$x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

where

- ▶ x is the original values of that feature for some data instance
 - ▶ x' is the scaled value of x , and
 - ▶ x_{\min} and x_{\max} are the lowest and highest values resp. in the dataset.
-
- ▶ Scaling maps the values of a feature into the interval $[0, 1]$

Feature Scaling (Min-Max)

Definition 2

Min-Max-Scaling a feature in a dataset D means computing

$$x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

where

- ▶ x is the original values of that feature for some data instance
 - ▶ x' is the scaled value of x , and
 - ▶ x_{\min} and x_{\max} are the lowest and highest values resp. in the dataset.
-
- ▶ Scaling maps the values of a feature into the interval $[0, 1]$
 - ▶ Scaling depends on the dataset. Changes yield differently scaled instances.

Feature Scaling (Min-Max)

Definition 2

Min-Max-Scaling a feature in a dataset D means computing

$$x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

where

- ▶ x is the original values of that feature for some data instance
 - ▶ x' is the scaled value of x , and
 - ▶ x_{\min} and x_{\max} are the lowest and highest values resp. in the dataset.
-
- ▶ Scaling maps the values of a feature into the interval $[0, 1]$
 - ▶ Scaling depends on the dataset. Changes yield differently scaled instances.
 - ▶ Scaling is invertible.

Outline

Basic Idea

The Neighborhood Size

The Decision Rule

Distance Functions

Python

kNN in Python

`sklearn.neighbors.KNeighborsClassifier`

Parameters

- ▶ k
- ▶ distance: large variety of out of the box functions
- ▶ decision rules: standard and weighted by distance
 - ▶ ties are resolved depending on the order of the data
- ▶ parameters controlling memory and cpu consumption

 Notebook 03_1_knn_digits

kNN in Python

`sklearn.neighbors.KNeighborsClassifier`

Parameters

- ▶ k
- ▶ distance: large variety of out of the box functions
- ▶ decision rules: standard and weighted by distance
 - ▶ ties are resolved depending on the order of the data
- ▶ parameters controlling memory and cpu consumption

 Notebook 03_1_knn_digits

 Exercises 1–3