

MADS-ML – Machine Learning

Ensemble Learning

Prof. Dr. Stephan Doerfel



FACHHOCHSCHULE KIEL
University of Applied Sciences



Moodle ()

Outline

Motivation

Voting

Bagging

Boosting

Why ensembles?

Why ensembles?

Idea Combine a set of weak learners into a stronger one.

Why ensembles?

- Idea** Combine a set of weak learners into a stronger one.
- ▶ Create a meta classifier.

Why ensembles?

- Idea** Combine a set of weak learners into a stronger one.
- ▶ Create a meta classifier.
 - ▶ Use a strategy to combine the predictions of the base learners into a single (final) classification.

Outline


Motivation

Voting

Bagging

Boosting

Voting Schemes



Source: [1], p. 220

Voting Schemes

In Python: `sklearn.ensemble.VotingClassifier`

Voting Schemes

In Python: `sklearn.ensemble.VotingClassifier`

Hard voting (plurality voting)

Voting Schemes

In Python: `sklearn.ensemble.VotingClassifier`

Hard voting (plurality voting)

- ▶ Each classifier votes for the class it predicted with weight 1

Voting Schemes

In Python: `sklearn.ensemble.VotingClassifier`

Hard voting (plurality voting)

- ▶ Each classifier votes for the class it predicted with weight 1
- ▶ The class with the most votes wins

Voting Schemes

In Python: `sklearn.ensemble.VotingClassifier`

Hard voting (plurality voting)

- ▶ Each classifier votes for the class it predicted with weight 1
- ▶ The class with the most votes wins
- ▶ Ties are broken implementation-specific, e.g. ordered by class label.

Voting Schemes

In Python: `sklearn.ensemble.VotingClassifier`

Hard voting (plurality voting)

- ▶ Each classifier votes for the class it predicted with weight 1
- ▶ The class with the most votes wins
- ▶ Ties are broken implementation-specific, e.g. ordered by class label.

Soft voting

Voting Schemes

In Python: `sklearn.ensemble.VotingClassifier`

Hard voting (plurality voting)

- ▶ Each classifier votes for the class it predicted with weight 1
- ▶ The class with the most votes wins
- ▶ Ties are broken implementation-specific, e.g. ordered by class label.

Soft voting

- ▶ Each classifier computes a probability for each class.

Voting Schemes

In Python: `sklearn.ensemble.VotingClassifier`

Hard voting (plurality voting)

- ▶ Each classifier votes for the class it predicted with weight 1
- ▶ The class with the most votes wins
- ▶ Ties are broken implementation-specific, e.g. ordered by class label.

Soft voting

- ▶ Each classifier computes a probability for each class.
- ▶ Per class, probabilities are summed up.

Voting Schemes

In Python: `sklearn.ensemble.VotingClassifier`

Hard voting (plurality voting)

- ▶ Each classifier votes for the class it predicted with weight 1
- ▶ The class with the most votes wins
- ▶ Ties are broken implementation-specific, e.g. ordered by class label.

Soft voting

- ▶ Each classifier computes a probability for each class.
- ▶ Per class, probabilities are summed up.
- ▶ Class with the highest sum wins.


The choice of the base predictors (types and number) and the voting scheme are new hyperparameters.

Rationale

Why is a combination of weak base learners more successful than a single one?

Rationale


Why is a combination of weak base learners more successful than a single one?



Source: [1], p. 223

Rationale

Why is a combination of weak base learners more successful than a single one?




Source: [1], p. 223

- idealized binary classification – **independent** error probabilities

Rationale

Why is a combination of weak base learners more successful than a single one?



Source: [1], p. 223

- ▶ idealized binary classification – **independent** error probabilities
- ▶ error probability (missclassification) of the ensemble is lower than that of any base classifier if the base classifiers is better than random guessing

Use Cases

- ▶ combine weak methods into a stronger one

Use Cases

- ▶ combine weak methods into a stronger one
- ▶ combine complementary approaches on the same data.

Use Cases

- ▶ combine weak methods into a stronger one
- ▶ combine complementary approaches on the same data.
- ▶ combine different data sources (each requiring different algorithms)

Outline

Motivation

Voting

Bagging

Boosting

Bagging

- ▶ use a voting scheme

Bagging

- ▶ use a voting scheme
- ▶ train each classifier on a different subset of the training data

Bagging

- ▶ use a voting scheme
- ▶ train each classifier on a different subset of the training data
- ▶ use random sampling with replacement to create the subsets.

Bagging

- ▶ use a voting scheme
- ▶ train each classifier on a different subset of the training data
- ▶ use random sampling with replacement to create the subsets.
- ▶ Python: `class sklearn.ensemble.BaggingClassifier`

Bagging

- ▶ use a voting scheme
- ▶ train each classifier on a different subset of the training data
- ▶ use random sampling with replacement to create the subsets.
- ▶ Python: `class sklearn.ensemble.BaggingClassifier`

 Notebook 07_1_ensembles_wines, Cells 1–7

Algorithm Random Forest

- ▶ **forest** = ensemble of trees
- ▶ bagging: each tree learned on a different data subset
- ▶ additionally: **feature bagging** each tree may only choose from a randomly sampled subset of features (random subspace) → more variation among the trees
- ▶ main hyperparameters: number of trees + hyperparameters of decision trees
- ▶ interpretability of decision trees is lost
- ▶ less prone to overfitting
- ▶ in sklearn: `sklearn.ensemble.RandomForestClassifier`

Outline

Motivation

Voting

Bagging

Boosting

Boosting

- ▶ train a sequence of classifiers

Boosting

- ▶ train a sequence of classifiers
- ▶ “focus” training of classifier i on missclassified instances of classifier $i - 1$

Boosting

- ▶ train a sequence of classifiers
- ▶ “focus” training of classifier i on missclassified instances of classifier $i - 1$
- ▶ warning: overfitting

Boosting

- ▶ train a sequence of classifiers
- ▶ “focus” training of classifier i on missclassified instances of classifier $i - 1$
- ▶ warning: overfitting
- ▶ most popular implementation: AdaBoost

Boosting

- ▶ train a sequence of classifiers
- ▶ “focus” training of classifier i on misclassified instances of classifier $i - 1$
- ▶ warning: overfitting
- ▶ most popular implementation: AdaBoost
- ▶ Python: `sklearn.ensemble.AdaBoostClassifier`

AdaBoost

- ▶ use a weight vector \mathbf{w} to control the influence of individual instances

AdaBoost

- ▶ use a weight vector \mathbf{w} to control the influence of individual instances
- ▶ \mathbf{w} has one entry for each sample

AdaBoost

- ▶ use a weight vector \mathbf{w} to control the influence of individual instances
- ▶ \mathbf{w} has one entry for each sample
- ▶ initialize \mathbf{w}^0 with ones

AdaBoost

- ▶ use a weight vector \mathbf{w} to control the influence of individual instances
- ▶ \mathbf{w} has one entry for each sample
- ▶ initialize \mathbf{w}^0 with ones
- ▶ for i in $1, 2, \dots, m$:

AdaBoost

- ▶ use a weight vector \mathbf{w} to control the influence of individual instances
- ▶ \mathbf{w} has one entry for each sample
- ▶ initialize \mathbf{w}^0 with ones
- ▶ for i in $1, 2, \dots, m$:
 - ▶ learn classifier C^i with training data weighted with \mathbf{w}^{i-1}

AdaBoost

- ▶ use a weight vector \mathbf{w} to control the influence of individual instances
- ▶ \mathbf{w} has one entry for each sample
- ▶ initialize \mathbf{w}^0 with ones
- ▶ for i in $1, 2, \dots, m$:
 - ▶ learn classifier C^i with training data weighted with \mathbf{w}^{i-1}
 - ▶ determine a coefficient based on the error rate of C^i

AdaBoost

- ▶ use a weight vector \mathbf{w} to control the influence of individual instances
- ▶ \mathbf{w} has one entry for each sample
- ▶ initialize \mathbf{w}^0 with ones
- ▶ for i in $1, 2, \dots, m$:
 - ▶ learn classifier C^i with training data weighted with \mathbf{w}^{i-1}
 - ▶ determine a coefficient based on the error rate of C^i
 - ▶ use the coefficient to increase weights of missclassified instances and decrease weights of correctly classified instances to yield \mathbf{w}^i

AdaBoost

- ▶ use a weight vector \mathbf{w} to control the influence of individual instances
- ▶ \mathbf{w} has one entry for each sample
- ▶ initialize \mathbf{w}^0 with ones
- ▶ for i in $1, 2, \dots, m$:
 - ▶ learn classifier C^i with training data weighted with \mathbf{w}^{i-1}
 - ▶ determine a coefficient based on the error rate of C^i
 - ▶ use the coefficient to increase weights of missclassified instances and decrease weights of correctly classified instances to yield \mathbf{w}^i
- ▶ Compute final prediction using a weighted combination of the m classifiers (weights are combinations of the respective coefficients and a learning rate)

AdaBoost – Algorithm

- ▶ For training data X, y , initialize $\mathbf{w}^{(0)}$ with $|X|$ ones and normalize

AdaBoost – Algorithm

- ▶ For training data X, y , initialize $\mathbf{w}^{(0)}$ with $|X|$ ones and normalize
- ▶ for i in $1, 2, \dots, m$:

AdaBoost – Algorithm

- ▶ For training data X, y , initialize $\mathbf{w}^{(0)}$ with $|X|$ ones and normalize
- ▶ for i in $1, 2, \dots, m$:
 - ▶ train $C^{(i)}$ on X, y with weights $\mathbf{w}^{(i-1)}$ and compute $\hat{y}^{(i)}$

AdaBoost – Algorithm

- ▶ For training data X, y , initialize $\mathbf{w}^{(0)}$ with $|X|$ ones and normalize
- ▶ for i in $1, 2, \dots, m$:
 - ▶ train $C^{(i)}$ on X, y with weights $\mathbf{w}^{(i-1)}$ and compute $\hat{y}^{(i)}$
 - ▶ error rate: $\epsilon^{(i)} := \sum_{x=1, y_x \neq \hat{y}_x^{(i)}}^{|X|} w_x^{(i-1)}$

AdaBoost – Algorithm

- ▶ For training data X, y , initialize $\mathbf{w}^{(0)}$ with $|X|$ ones and normalize
- ▶ for i in $1, 2, \dots, m$:
 - ▶ train $C^{(i)}$ on X, y with weights $\mathbf{w}^{(i-1)}$ and compute $\hat{y}^{(i)}$
 - ▶ error rate: $\epsilon^{(i)} := \sum_{x=1, y_x \neq \hat{y}_x^{(i)}}^{|X|} w_x^{(i-1)}$
 - ▶ coefficient $\alpha^{(i)} := 0.5 \log \frac{1-\epsilon^{(i)}}{\epsilon^{(i)}}$

AdaBoost – Algorithm

- ▶ For training data X, y , initialize $\mathbf{w}^{(0)}$ with $|X|$ ones and normalize
- ▶ for i in $1, 2, \dots, m$:
 - ▶ train $C^{(i)}$ on X, y with weights $\mathbf{w}^{(i-1)}$ and compute $\hat{y}^{(i)}$
 - ▶ error rate: $\epsilon^{(i)} := \sum_{x=1, y_x \neq \hat{y}_x^{(i)}}^{|X|} w_x^{(i-1)}$
 - ▶ coefficient $\alpha^{(i)} := 0.5 \log \frac{1-\epsilon^{(i)}}{\epsilon^{(i)}}$
 - ▶ update weights $w_x^{(i)} := w_x^{(i-1)} e^{-y_x \hat{y}_x^{(i)} \alpha^{(i)}}$ and normalize $\mathbf{w}^{(i)}$

AdaBoost – Algorithm

- ▶ For training data X, y , initialize $\mathbf{w}^{(0)}$ with $|X|$ ones and normalize
- ▶ for i in $1, 2, \dots, m$:
 - ▶ train $C^{(i)}$ on X, y with weights $\mathbf{w}^{(i-1)}$ and compute $\hat{y}^{(i)}$
 - ▶ error rate: $\epsilon^{(i)} := \sum_{x=1, y_x \neq \hat{y}_x^{(i)}}^{|X|} w_x^{(i-1)}$
 - ▶ coefficient $\alpha^{(i)} := 0.5 \log \frac{1-\epsilon^{(i)}}{\epsilon^{(i)}}$
 - ▶ update weights $w_x^{(i)} := w_x^{(i-1)} e^{-y_x \hat{y}_x^{(i)} \alpha^{(i)}}$ and normalize $\mathbf{w}^{(i)}$

Note: If $C^{(i)}$ is better than guessing, then $\epsilon^{(i)} < 0.5$. Thus,

$$\alpha^{(i)} = .5 \log \frac{1-\epsilon^{(i)}}{\epsilon^{(i)}} > 0 \text{ and } e^{-y_x \hat{y}_x^{(i)} \alpha^{(i)}} \begin{cases} > 1 \text{ for } y_x \neq \hat{y}_x^{(i)} \\ < 1 \text{ for } y_x = \hat{y}_x^{(i)} \end{cases}.$$

AdaBoost in Python

► `csklearn.ensemble.AdaBoostClassifier`

AdaBoost in Python

- ▶ `csklearn.ensemble.AdaBoostClassifier`
- ▶ parameters: estimator type, number of estimators

AdaBoost in Python

- ▶ `csklearn.ensemble.AdaBoostClassifier`
- ▶ parameters: estimator type, number of estimators
- ▶ additionally use a learning rate to control how much the i -th learner contributes to the result of the previous $i - 1$ learners

AdaBoost in Python

- ▶ `csklearn.ensemble.AdaBoostClassifier`
- ▶ parameters: estimator type, number of estimators
- ▶ additionally use a learning rate to control how much the i -th learner contributes to the result of the previous $i - 1$ learners

 Notebook 07_1_ensembles_wines, Cells 8–10

References



S. Raschka and V. Mirjalili.

Python Machine Learning.

Packt Publishing Ltd., Livery Place 35 Livery Street Birmingham B3 2PB, UK, second edition, September 2017.